

MITHUN REDDY, SUDIREDDY

📍 Hyderabad, India 📞 +91 9676078758 ✉ mithunreddy@outlook.in
🌐 <https://mithunr.app/> 📺 mithunreddyyy 🌱 mithunreddyy

PROFESSIONAL SUMMARY

Full-Stack Software Engineer & AI Enthusiast with over 2+ years of experience in designing scalable web architectures and machine learning pipelines. Proven ability to optimize backend databases, build retrieval-augmented generation systems, and automate data workflows to improve system performance. Skilled in React, Node.js, Python, and cloud platforms (AWS, GCP).

EDUCATION

ICFAI University - Hyderabad, India	Aug 2017 - Jul 2021
Bachelor of Technology in Computer Science Engineering	
Coursework: Data Structures and Algorithms, Software engineering, Probability, Statistics, Artificial Intelligence, Machine Learning, Computer Networking, Databases, Distributed Systems, Cloud Computing, Operating Systems	

SKILLS

- Programming Languages:** TypeScript, JavaScript (ES6+), Python, SQL
- AI/ML Engineering:** LangChain, LangGraph, HuggingFace, RAG, Vector Databases (FAISS, ChromaDB, Qdrant), Transformers, PyTorch, TensorFlow, scikit-learn, XGBoost, Pandas, NumPy, MLOps, MLflow
- Big Data & Data Engineering:** Apache Spark / PySpark, Databricks, SQL, PostgreSQL, MongoDB, Redis, Cassandra
- Cloud Platforms:** GCP (Cloud Run, Cloud SQL, Cloud Storage, Cloud Scheduler, IAM, VPC networks), AWS (EC2, S3, EKS, SageMaker), Azure (Container Apps, App Service, Storage, Key Vault)
- Infrastructure & DevOps:** FastAPI, GitHub Actions (CI/CD), Docker, TerraForm
- Teamwork:** Git, GitHub, Jira, Confluence, Agile/Scrum
- Certifications:** Associate Cloud Engineer - Google Cloud

EXPERIENCE

Full-Stack Engineer: Pentagonam InfoTech (Hyderabad, India)	Feb 2024 - Present
- Architected and delivered a modular microservices ecosystem supporting authentication, order management, analytics, reporting, and workflow automation, reducing system dependencies and improving overall platform maintainability by 45%. - Built highly optimized RESTful APIs and backend services handling millions of monthly transactions while maintaining sub-200ms average response times through database tuning, indexing strategies, caching layers, and asynchronous processing. - Led a team of 6 developers across Frontend, Backend, and DevOps functions, conducting sprint planning, architecture reviews, code reviews, technical mentoring, and release management within Agile delivery environments. - Established engineering standards, reusable component libraries, coding guidelines, and CI/CD workflows, increasing development velocity by 35% while reducing production defects. - Developed advanced data processing pipelines using Python, SQL, and PySpark to automate business intelligence workflows, reducing manual operational effort by more than 20 hours per week. - Optimized PostgreSQL database performance through partitioning, query plan optimization, materialized views, connection pooling, and indexing strategies, improving critical transaction throughput by over 50%. - Implemented Infrastructure as Code using Terraform and Docker-based deployment pipelines integrated with GitHub Actions, enabling fully automated build, testing, security scanning, and production deployments. - Successfully delivered 20+ major product releases with zero critical production incidents, maintaining 99.95% platform availability and supporting continuous business growth.	
Software Engineer: Zensar Technologies (Hyderabad, India)	Jan 2022 - Nov 2023
- Led the transformation of legacy monolithic services into scalable event-driven microservice architectures using Node.js, Express.js, Python, AWS SQS, and SNS, enabling resilient processing of high-volume enterprise transactions with near-zero message failures. - Re-engineered PostgreSQL database infrastructure through advanced indexing strategies, query plan optimization, partitioning, materialized views, and connection pooling, reducing API response times from over 800ms to below 200ms across critical business workflows. - Designed distributed caching and data acceleration layers using Redis and AWS CloudFront, decreasing backend database dependency by 60% while enabling real-time analytics dashboards with sub-100ms response performance. - Built cloud-native infrastructure across AWS services including EC2, S3, RDS, Lambda, and CloudFront using Terraform-based Infrastructure as Code practices, achieving 99.99% platform availability and streamlined environment provisioning. - Established fully automated CI/CD ecosystems using Docker, GitHub Actions, automated testing pipelines, and deployment orchestration, reducing release cycles by 70% while enabling zero-downtime production deployments.	

- Engineered secure authentication and authorization frameworks leveraging OAuth 2.0, RBAC, API throttling, and enterprise security controls to safeguard applications operating within highly regulated financial and healthcare environments.
- Collaborated directly with solution architects, product owners, QA teams, business stakeholders, and offshore delivery teams to drive end-to-end software delivery, technical planning, sprint execution, and production support initiatives.
- Mentored junior developers through code reviews, architectural discussions, debugging sessions, and engineering best practices while contributing to improved development standards, release quality, and team productivity.

PROJECTS

Autonomous Career Architect [Local Job Application Pipeline with BM25 and LaTeX]

- (Python, BM25, LaTeX, Streamlit, Ollama, Pandoc, JSON)
- Built a zero-cost CLI pipeline that parses job descriptions into structured role signals, ranks resume and project evidence using BM25, and generates tailored resume, cover letter, interview prep, and ATS report artifacts per application with no paid API dependencies.
 - Rendered output into four LaTeX templates (default, minimal, creative, executive) with ATS keyword alignment scoring, version snapshots with diff support, batch job processing, HTML and remote URL scraping, and a Streamlit dashboard for reviewing application stats.

LocalPilot [Private Code Completion with Local LLM Inference]

- (Python, TypeScript, FastAPI, Ollama, CodeLlama, Qdrant, BM25, tree-sitter, SSE)
- Built a fully local VS Code extension and FastAPI backend delivering inline completions, code explanations, test generation, and refactoring via Server-Sent Events with 400ms debounced streaming. Ran CodeLlama 13B or 7B through Ollama with zero cloud dependencies and no API keys.
 - Implemented hybrid codebase indexing using Qdrant vector search with 768-dimensional nomic-embed-code embeddings and BM25 keyword search over Python ASTs parsed with tree-sitter. Combined semantic and symbol-level retrieval to supply relevant context from the full workspace into every completion request.

PromptVault [Version-Controlled Prompt Management with Evaluation Tracking]

- (Python, FastAPI, SQLite, React, TypeScript, Ollama, Docker)
- Built a local prompt management system using FastAPI and SQLite. The application stores prompt versions with full metadata for model parameters and system instructions. The backend supports tag-based organization and side-by-side version comparisons. The API provides endpoints for creation, rollback, and search operations.
 - Implemented a structured evaluation layer to test prompt variants against fixed datasets. The pipeline scores outputs using exact match algorithms or a local Ollama model. A React dashboard displays win-rate comparisons, score distributions, and regression alerts. The system exports evaluation reports in JSON and CSV formats.

INTERESTS

Data-driven Problem Solving, Exploring New Technologies, Reading, Video Editing, Motion Design, Music, Watching Movies.

ADDITIONAL INFORMATION

- **Languages:** Telugu (Native), Hindi (Fluent), English (Proficient)
- **Willingness to Relocate:** Open to relocation within remote opportunities