

MITHUN REDDY, SUDIREDDY

📍 Hyderabad, India 📞 +91 9676078758 ✉ mithunreddy@outlook.in

🌐 <https://mithunr.app/> 📺 mithunreddyyy 🌐 mithunreddyy

PROFESSIONAL SUMMARY

Full-Stack Software Engineer & AI Enthusiast with over 2+ years of experience in designing scalable web architectures and machine learning pipelines. Proven ability to optimize backend databases, build retrieval-augmented generation systems, and automate data workflows to improve system performance. Skilled in React, Node.js, Python, local LLM inference, and cloud platforms (AWS, GCP).

- Full-Stack Engineering / React.js / Node.js
- Python / TypeScript / PostgreSQL / PySpark
- AI Engineering / RAG / Agentic AI
- AWS / GCP / Azure / CI/CD

EDUCATION

ICFAI University - Hyderabad, India

Aug 2017 - Jul 2021

Bachelor of Technology in Computer Science Engineering

- Coursework: Data Structures and Algorithms, Software engineering, Probability, Statistics, Artificial Intelligence, Machine Learning, Computer Networking, Databases, Distributed Systems, Cloud Computing, Operating Systems

SKILLS

- **Programming Languages:** TypeScript, JavaScript (ES6+), Python, SQL, HTML5, CSS3
- **AI/ML Engineering:** LangChain, LangGraph, HuggingFace, RAG, Vector Databases (FAISS, ChromaDB, Qdrant), Fine-tuning, PyTorch, TensorFlow, scikit-learn, XGBoost, Pandas, NumPy, MLOps, MLflow, W&B
- **Big Data & Data Engineering:** Apache Spark / PySpark, Databricks, SQL, PostgreSQL, MongoDB, Redis, Cassandra
- **Cloud Platforms:** GCP (Cloud Run, Cloud SQL, Cloud Storage, Cloud Scheduler, IAM, VPC networks), AWS (EC2, S3, EKS, SageMaker), Azure (Container Apps, App Service, Storage, Key Vault)
- **Infrastructure & DevOps:** FastAPI, GitHub Actions (CI/CD), Docker, TerraForm
- **Teamwork:** Git, GitHub, Jira, Confluence, Agile/Scrum
- **Certifications:** Associate Cloud Engineer - Google Cloud

EXPERIENCE

Full-Stack Engineer: Zensar Technologies (Hyderabad, India)

Jan 2022 - Nov 2023

- Engineered high-performance web architectures using React.js, TypeScript, and Next.js. Implemented Server-Side Rendering (SSR) and advanced state management to improve Core Web Vitals and reduce Time to Interactive (TTI) by 40%.
- Architected scalable backend microservices using Node.js, Express, and Python. Transitioned monolithic endpoints to an event-driven architecture using AWS SQS/SNS to process high-throughput concurrent enterprise traffic without message loss.
- Optimized PostgreSQL database operations through table partitioning, B-tree/GIN indexing, and connection pooling. Restructured complex query execution plans to decrease average REST API response latency from 800 ms to under 200 ms.
- Designed distributed caching layers using Redis and AWS CloudFront. Decreased direct database read loads by 60% and enabled sub-100ms data retrieval for mission-critical real-time client dashboards.
- Provisioned fault-tolerant cloud infrastructure across AWS (EC2, S3, RDS, Lambda) using Terraform for Infrastructure as Code (IaC). Established highly available production environments ensuring 99.99% system uptime.
- Automated CI/CD pipelines using GitHub Actions and Docker. Containerized application deployments to establish consistent staging environments, achieving zero-downtime rolling updates and reducing overall deployment times by 70%.
- Developed automated, fault-tolerant ETL data pipelines using Python, SQL, and PySpark to process high-volume enterprise data. Eliminated 15 hours of manual operational overhead per week while maintaining strict ACID compliance and 99.9% data integrity.
- Implemented enterprise security protocols including OAuth 2.0 token-based authentication, Role-Based Access Control (RBAC), and strict API rate-limiting to secure systems within the highly regulated FinTech and Healthcare domains.

Frontend Engineer Intern: Sagiam (Hyderabad, India)

Aug 2021 - Jan 2022

- Architected reusable, type-safe UI component libraries using React and TypeScript for internal dashboards and customer-facing portals. Implemented Redux for strict state management across multiple concurrent data streams.
- Optimized client-side rendering performance by implementing code splitting, lazy loading, and component memoization. Improved initial page load speed by 30% and achieved a 90+ Lighthouse performance score.
- Developed backend API endpoints using Node.js to aggregate CRM and financial data. Served optimized JSON payloads to the frontend, reducing client-side data processing overhead and improving interface responsiveness.

PROJECTS

Autonomous Career Architect [Local Job Application Pipeline with BM25 and LaTeX]

- (Python, BM25, LaTeX, Streamlit, Ollama, Pandoc, JSON)
- Built a zero-cost CLI pipeline that parses job descriptions into structured role signals, ranks resume and project evidence using BM25, and generates tailored resume, cover letter, interview prep, and ATS report artifacts per application with no paid API dependencies.
 - Rendered output into four LaTeX templates (default, minimal, creative, executive) with ATS keyword alignment scoring, version snapshots with diff support, batch job processing, HTML and remote URL scraping, and a Streamlit dashboard for reviewing application stats.

LocalPilot [Private Code Completion with Local LLM Inference]

- (Python, TypeScript, FastAPI, Ollama, CodeLlama, Qdrant, BM25, tree-sitter, SSE)
- Built a fully local VS Code extension and FastAPI backend delivering inline completions, code explanations, test generation, and refactoring via Server-Sent Events with 400ms debounced streaming. Ran CodeLlama 13B or 7B through Ollama with zero cloud dependencies and no API keys.
 - Implemented hybrid codebase indexing using Qdrant vector search with 768-dimensional nomic-embed-code embeddings and BM25 keyword search over Python ASTs parsed with tree-sitter. Combined semantic and symbol-level retrieval to supply relevant context from the full workspace into every completion request.

PromptVault [Version-Controlled Prompt Management with Evaluation Tracking]

- (Python, FastAPI, SQLite, React, TypeScript, Ollama, Docker)
- Built a local prompt management system using FastAPI and SQLite. The application stores prompt versions with full metadata for model parameters and system instructions. The backend supports tag-based organization and side-by-side version comparisons. The API provides endpoints for creation, rollback, and search operations.
 - Implemented a structured evaluation layer to test prompt variants against fixed datasets. The pipeline scores outputs using exact match algorithms or a local Ollama model. A React dashboard displays win-rate comparisons, score distributions, and regression alerts. The system exports evaluation reports in JSON and CSV formats.

INTERESTS

Data-driven Problem Solving, Exploring New Technologies, AI & Automation, Reading, Video Editing, Motion Design, Music, Watching Movies.

ADDITIONAL INFORMATION

- **Languages:** Telugu (Native), Hindi (Fluent), English (Proficient)
- **Willingness to Relocate:** Open to relocation within remote opportunities